

Machine Learning for Bitcoin Price Prediction and Sentiment Analysis

In this project I have two problems I want to solve, they are

1. I want to recognize patterns in BTC/USDT to forecast what will happen to the price in the next 12 hours with some technical indicators (Engineering features)
2. I want to get sentiment analysis from any source like CoinDesk or CryptoPanic

Let's focus on number one first, right now, we need bitcoin price data, it's more important than anything else, but wait... how can we get up-to-date data from the market? Ohh I know, previously I had a similar project like this and I used ccxt to get bitget data. Oke, lets get bitcoin data from bitget...

DISCLAIMER

Everything I do here is just for brain training and personal experimentation. Not everything should be trusted or used as a benchmark for investment or trading. The goal is just to learn and apply the Data Science knowledge that I have learned so far.

```
In [1]: # pip install pycoingecko ccxt ta sncscrape
```

```
In [2]: import pandas as pd
import ccxt
```

Get Data from API

```
In [3]: bitget = ccxt.bitget()
ohlcv = bitget.fetch_ohlcv('BTC/USDT', timeframe='1h', limit=1000)
```

EDA (Exploratory Data Analysis)

```
In [4]: df = pd.DataFrame(ohlcv, columns=['timestamp', 'open', 'high', 'low', 'close', 'volume'])
df['timestamp'] = pd.to_datetime(df['timestamp'], unit='ms')
df.set_index('timestamp', inplace=True)

df
```

```
Out[4]:
```

	open	high	low	close	volume
timestamp					
2025-03-04 15:00:00	83216.68	83387.09	81520.90	82999.99	1002.587290
2025-03-04 16:00:00	82999.99	83790.00	82022.41	83552.68	1179.901143
2025-03-04 17:00:00	83552.68	85530.27	83429.00	85450.72	884.769535
2025-03-04 18:00:00	85450.72	86900.00	85173.18	86765.01	812.538519
2025-03-04 19:00:00	86765.01	88467.73	86535.02	88200.01	1195.767062
...
2025-04-15 02:00:00	84966.20	85284.50	84697.47	85220.32	296.418309
2025-04-15 03:00:00	85220.32	85403.63	85068.00	85365.43	345.552030
2025-04-15 04:00:00	85365.43	85511.68	85155.07	85485.87	314.500596
2025-04-15 05:00:00	85485.87	85855.86	85473.01	85535.48	373.851146
2025-04-15 06:00:00	85535.48	85678.00	85362.00	85400.22	217.607356

1000 rows × 5 columns

Right now we have the data, but we don't know if the data is good or there are actually zero values or missing values? But I'm pretty sure if the data is up-to-date from the market, we won't have zero/missing values, but let's check first

```
In [5]: #Check Missing Values
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1000 entries, 2025-03-04 15:00:00 to 2025-04-15 06:00:00
Data columns (total 5 columns):
#   Column   Non-Null Count  Dtype
---  -
0    open     1000 non-null   float64
1    high     1000 non-null   float64
2    low      1000 non-null   float64
3    close    1000 non-null   float64
4    volume   1000 non-null   float64
dtypes: float64(5)
memory usage: 46.9 KB
```

```
Out[5]:
```


	open	high	low	close	volume
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	83897.264280	84246.455350	83561.982620	83899.447820	375.060114
std	2808.649191	2784.765309	2849.135769	2808.968313	432.975775
min	74828.000000	75486.720000	74522.270000	74828.000000	24.073424
25%	82567.252500	82881.307500	82249.995000	82567.252500	149.520707
50%	83883.450000	84126.735000	83617.345000	83889.445000	244.524613
75%	85345.312500	85716.702500	85093.030000	85356.940000	433.774531
max	92728.210000	92794.110000	92135.000000	92728.210000	5328.027275

okay, thats good, we dont have missing values, lets see the price of bitcoin in the last 1000 candles

```
In [6]: import matplotlib.pyplot as plt
import matplotlib.dates as mdates
```

```
In [7]: #Check Linechar, i want to see the price of BTC in the last 1000 candles
```

```
plt.plot(df.index, df['close'], label='BTC Price')
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%m-%d'))
plt.xlabel('Time')
plt.ylabel('Price')
plt.title('BTC Price in the Last 1000 Candles')
plt.legend()
plt.show()
```



Analyze Data

Previously we see the data and i think the data is healthy, we can up the stairs. In this part we will analyze the data, and we have to recognize what features make a good prediction. Actually, I don't have the experience to know what features make a good prediction. I am a trader and I know that the market is actually unpredictable and actually traders only use historical data to predict the future... and that's hard.

But, lets get the technical indicator first, i will use RSI, ATR, Bollinger Bands, Volatility, etc. (We will use ta library)

```
In [8]: import numpy as np
from ta.momentum import RSIIndicator
from ta.volatility import BollingerBands
from ta.volatility import AverageTrueRange

atr = AverageTrueRange(df['high'], df['low'], df['close'], window=14)
bb = BollingerBands(close=df['close'], window=20, window_dev=2)
```

```
In [9]: df_analyze = df.copy() # Copy data

df_analyze['close_future_6'] = df_analyze['close'].shift(6)
df_analyze['price_movement'] = (df_analyze['close_future_6'] > df_analyze['close']).astype(int)
df_analyze['close_future_1'] = df_analyze['close'].shift(3)
df_analyze['candle_body_pct'] = (df_analyze['close'] - df_analyze['open']) / (df_analyze['high'] - df_analyze['low'])
df_analyze['volume_change'] = df_analyze['volume'].diff()
df_analyze['bb_width'] = bb.bollinger_hband() - bb.bollinger_lband()
df_analyze['close_pct_1'] = df_analyze['close'].pct_change(1)
df_analyze['return_lag_1'] = df_analyze['close_pct_1'].shift(1)
df_analyze['volatility_5'] = df_analyze['close_pct_1'].rolling(5).std()
df_analyze['ema_10'] = df_analyze['close'].ewm(span=10).mean()
df_analyze['ema_20'] = df_analyze['close'].ewm(span=20).mean()
df_analyze['ema_10_rel'] = df_analyze['close'] / df_analyze['close'].ewm(span=10).mean() - 1
df_analyze['ema_20_rel'] = df_analyze['close'] / df_analyze['close'].ewm(span=20).mean() - 1
df_analyze['ema_cross'] = ((df_analyze['ema_10'] > df_analyze['ema_20']) &
                           (df_analyze['ema_10'].shift(1) <= df_analyze['ema_20'].shift(1))).astype(int)

df_analyze['rsi'] = RSIIndicator(df_analyze['close'], window=14).rsi()
df_analyze['rsi_change'] = df_analyze['rsi'] - df_analyze['rsi'].shift(6)
df_analyze['bb_width_change'] = df_analyze['bb_width'].pct_change(3)
df_analyze['vol_accel'] = df_analyze['volume'].pct_change().pct_change()
df_analyze['returns'] = df_analyze['close'].pct_change() * 100
df_analyze['atr'] = atr.average_true_range()
df_analyze['wick_upper'] = df[['high', 'close']].max(axis=1)
df_analyze['wick_lower'] = df[['open', 'close']].min(axis=1) - df['low']
df_analyze['price_activity_ratio'] = (df_analyze['high'] - df_analyze['low']) / (df_analyze['high'].rolling(12).mean() - df_analyze['low'].rolling(12).mean())
df_analyze['volume_ratio_4h'] = df_analyze['volume'] / df_analyze['volume'].rolling(4).mean()
df_analyze['rsi_slope'] = df_analyze['rsi'] - df_analyze['rsi'].shift(3)
df_analyze['atr_change'] = df_analyze['atr'].pct_change(4)
df_analyze['close_momentum_6h'] = df_analyze['close'].shift(6) - 1
df_analyze['close_momentum_12h'] = df_analyze['close'].shift(12) - 1
df_analyze['close_momentum_24h'] = df_analyze['close'].shift(24) - 1
df_analyze['price_zscore'] = (df_analyze['close'] - df_analyze['close'].rolling(50).mean()) / df_analyze['close'].rolling(50).std()
df_analyze['hour'] = df_analyze.index.hour
df_analyze['hour_sin'] = np.sin(2 * np.pi * df_analyze.index.hour / 24)
df_analyze['hour_cos'] = np.cos(2 * np.pi * df_analyze.index.hour / 24)
df_analyze['day_of_week'] = df_analyze.index.dayofweek
df_analyze['close_pct_3'] = df_analyze['close'].pct_change(3)
df_analyze['log_return'] = np.log(df_analyze['close'] / df_analyze['close'].shift(1))
df_analyze['close_lag_1'] = df_analyze['close'].shift(1)
df_analyze['hl_range'] = df_analyze['high'] - df_analyze['low']

df_analyze.tail()
```

```
Out[9]:
```

	open	high	low	close	volume	close_future_6	price_movement	close_future_1	candle_body_pct
timestamp									
2025-04-15 02:00:00	84966.20	85284.50	84697.47	85220.32	296.418309	84860.07	0	84599.99	0.432891
2025-04-15 03:00:00	85220.32	85403.63	85068.00	85365.43	345.552030	84700.03	0	84934.24	0.432351
2025-04-15 04:00:00	85365.43	85511.68	85155.07	85485.87	314.500596	84517.01	0	84966.20	0.337736
2025-04-15 05:00:00	85485.87	85855.86	85473.01	85535.48	373.851146	84599.99	0	85220.32	0.129581
2025-04-15 06:00:00	85535.48	85678.00	85362.00	85400.22	217.607356	84934.24	0	85365.43	-0.428038

5 rows × 43 columns

```
In [10]: print("Sebelum dilakukan handling")
print(df_analyze.info())

df_analyze.dropna(inplace=True)
print("\n-----\n")
```

```
print("Setelah dilakukan handling")
print(df_analyze.info())
```

Sebelum dilakukan handling

```
<class 'pandas.core.frame.DataFrame'>
```

DatetimeIndex: 1000 entries, 2025-03-04 15:00:00 to 2025-04-15 06:00:00

Data columns (total 43 columns):

#	Column	Non-Null Count	Dtype
0	open	1000 non-null	float64
1	high	1000 non-null	float64
2	low	1000 non-null	float64
3	close	1000 non-null	float64
4	volume	1000 non-null	float64
5	close_future_6	994 non-null	float64
6	price_movement	1000 non-null	int64
7	close_future_1	997 non-null	float64
8	candle_body_pct	1000 non-null	float64
9	volume_change	999 non-null	float64
10	bb_width	981 non-null	float64
11	close_pct_1	999 non-null	float64
12	return_lag_1	998 non-null	float64
13	volatility_5	995 non-null	float64
14	ema_10	1000 non-null	float64
15	ema_20	1000 non-null	float64
16	ema_10_rel	1000 non-null	float64
17	ema_20_rel	1000 non-null	float64
18	ema_cross	1000 non-null	int64
19	rsi	987 non-null	float64
20	rsi_change	981 non-null	float64
21	bb_width_change	978 non-null	float64
22	vol_accel	998 non-null	float64
23	returns	999 non-null	float64
24	atr	1000 non-null	float64
25	wick_upper	1000 non-null	float64
26	wick_lower	1000 non-null	float64
27	price_activity_ratio	989 non-null	float64
28	volume_ratio_4h	997 non-null	float64
29	rsi_slope	984 non-null	float64
30	atr_change	987 non-null	float64
31	close_momentum_6h	994 non-null	float64
32	close_momentum_12h	988 non-null	float64
33	close_momentum_24h	976 non-null	float64
34	price_zscore	951 non-null	float64
35	hour	1000 non-null	int32
36	hour_sin	1000 non-null	float64
37	hour_cos	1000 non-null	float64
38	day_of_week	1000 non-null	int32
39	close_pct_3	997 non-null	float64
40	log_return	999 non-null	float64
41	close_lag_1	999 non-null	float64
42	hl_range	1000 non-null	float64

dtypes: float64(39), int32(2), int64(2)

memory usage: 335.9 KB

None

Setelah dilakukan handling

```
<class 'pandas.core.frame.DataFrame'>
```

DatetimeIndex: 951 entries, 2025-03-06 16:00:00 to 2025-04-15 06:00:00

Data columns (total 43 columns):

#	Column	Non-Null Count	Dtype
0	open	951 non-null	float64
1	high	951 non-null	float64
2	low	951 non-null	float64
3	close	951 non-null	float64
4	volume	951 non-null	float64
5	close_future_6	951 non-null	float64
6	price_movement	951 non-null	int64
7	close_future_1	951 non-null	float64
8	candle_body_pct	951 non-null	float64
9	volume_change	951 non-null	float64
10	bb_width	951 non-null	float64
11	close_pct_1	951 non-null	float64
12	return_lag_1	951 non-null	float64
13	volatility_5	951 non-null	float64
14	ema_10	951 non-null	float64
15	ema_20	951 non-null	float64
16	ema_10_rel	951 non-null	float64
17	ema_20_rel	951 non-null	float64
18	ema_cross	951 non-null	int64
19	rsi	951 non-null	float64

```

20 rsi_change          951 non-null    float64
21 bb_width_change     951 non-null    float64
22 vol_accel           951 non-null    float64
23 returns             951 non-null    float64
24 atr                 951 non-null    float64
25 wick_upper          951 non-null    float64
26 wick_lower          951 non-null    float64
27 price_activity_ratio 951 non-null    float64
28 volume_ratio_4h     951 non-null    float64
29 rsi_slope           951 non-null    float64
30 atr_change          951 non-null    float64
31 close_momentum_6h   951 non-null    float64
32 close_momentum_12h  951 non-null    float64
33 close_momentum_24h  951 non-null    float64
34 price_zscore        951 non-null    float64
35 hour                951 non-null    int32
36 hour_sin            951 non-null    float64
37 hour_cos            951 non-null    float64
38 day_of_week         951 non-null    int32
39 close_pct_3         951 non-null    float64
40 log_return          951 non-null    float64
41 close_lag_1         951 non-null    float64
42 hl_range            951 non-null    float64
dtypes: float64(39), int32(2), int64(2)
memory usage: 319.5 KB
None

```

That's... quite a lot isn't it hehehe? honestly, I'm still learning and I wasn't sure which features would give the best prediction results. So I decided to try as many reasonable features as I could think of (mostly from technical indicators and price-related patterns). Later on, I can use feature selection or model evaluation to figure out which ones actually help. For now, it's more like learning by doing and seeing what works.

Model Machine Learning

You all know about LSTM or Long Short Term Memory? i will use that in this prediction, reason? i dont know, ups sorry, i know the reason hehehe. Because I think instead of using linear regression because the current condition of bitcoin is very volatile and linear regression will be overwhelmed in predicting it, therefore I choose LSTM because it can handle volatility (but I also know that LSTM has the potential for overfitting especially the features I use a lot lol).

For this project, I'm predicting two targets: the current price, and the price from six hours ago (by shifting the price column upward). My idea is to give the model a short-term and slightly mid-term perspective, so it can learn both recent patterns and a bit of context from the past.

```

In [11]: from sklearn.preprocessing import MinMaxScaler
from keras.models import Model
from keras.layers import Input, LSTM, Dense, Dropout
from keras.callbacks import EarlyStopping
from sklearn.model_selection import TimeSeriesSplit
from sklearn.feature_selection import SelectKBest, f_regression

```

```

In [12]: selected_features = [
    'open', 'high', 'low', 'volume',
    'candle_body_pct', 'volume_change', 'bb_width', 'close_pct_1', 'return_lag_1',
    'volatility_5', 'ema_10_rel', 'ema_20_rel',
    'rsi', 'rsi_change', 'rsi_slope',
    'bb_width_change', 'vol_accel', 'returns', 'atr', 'atr_change',
    'wick_upper', 'wick_lower', 'price_activity_ratio', 'volume_ratio_4h',
    'close_momentum_6h', 'close_momentum_12h', 'close_momentum_24h',
    'price_zscore', 'log_return', 'hl_range'
]

selector = SelectKBest(f_regression, k=20)
X_selected = selector.fit_transform(df_analyze[selected_features], df_analyze['close'])

selected_features = [selected_features[i] for i in selector.get_support(indices=True)]
print("Selected features:", selected_features)

y_multi = df_analyze[['close', 'close_future_6']].dropna()

y_data_reg = df_analyze['close'].values.reshape(-1, 1)
y_data_clf = df_analyze['price_movement'].values

scaler_X = MinMaxScaler()
scaler_y = MinMaxScaler()

X_aligned = df_analyze.loc[y_multi.index, selected_features].values

X_scaled = scaler_X.fit_transform(X_aligned)
y_scaled_reg = scaler_y.fit_transform(y_data_reg)
y_scaled_multi = scaler_y.fit_transform(y_multi.values)

```

```
df_scaled = np.hstack([y_scaled_multi, X_scaled])
```

```
df_scaled
```

Selected features: ['open', 'high', 'low', 'volume', 'bb_width', 'close_pct_1', 'volatility_5', 'ema_10_rel', 'ema_20_rel', 'rsi', 'atr', 'atr_change', 'wick_upper', 'wick_lower', 'close_momentum_6h', 'close_momentum_12h', 'close_momentum_24h', 'price_zscore', 'log_return', 'hl_range']

```
Out[12]: array([[0.92862374, 0.99635296, 1.          , ..., 0.58766014, 0.33924641,
                  0.27392325],
                [0.88456255, 1.          , 0.90668462, ..., 0.53909206, 0.42829386,
                  0.20237884],
                [0.87106632, 0.92463536, 0.8636644 , ..., 0.51458006, 0.48222187,
                  0.11466484],
                ...,
                [0.68222856, 0.58639533, 0.6585832 , ..., 0.79334915, 0.52041083,
                  0.04984241],
                [0.68540418, 0.59141742, 0.66611063, ..., 0.80324004, 0.5120529 ,
                  0.05386488],
                [0.67674595, 0.6116468 , 0.66921123, ..., 0.76191743, 0.49026712,
                  0.04361709]])
```

I split the data manually, I don't use train-test-split because it will make the prediction overfitting (because the past will learn the future, it's crazy, it doesn't make sense). By the way, I split the data into two parts, first regression and classification, I just want the model to have more information about the price (and to understand about volatility).

```
In [13]: window_size = 60

def create_sequences(X, y, window=60):
    Xs, ys = [], []
    for i in range(window, len(X)):
        Xs.append(X[i - window:i])
        ys.append(y[i])
    return np.array(Xs), np.array(ys)

X_seq, y_seq_reg = create_sequences(X_scaled, y_scaled_reg, window=window_size)
_, y_seq_clf = create_sequences(X_scaled, y_data_clf, window=window_size)

X_seq, y_seq = create_sequences(X_scaled, y_scaled_multi, window_size)

split_index = int(len(X_seq) * 0.7)
X_train, X_test = X_seq[:split_index], X_seq[split_index:]
y_train_reg, y_test_reg = y_seq_reg[:split_index], y_seq_reg[split_index:]
y_train_clf, y_test_clf = y_seq_clf[:split_index], y_seq_clf[split_index:]
```

```
In [14]: input_layer = Input(shape=(X_train.shape[1], X_train.shape[2]))
x = LSTM(128, return_sequences=True)(input_layer)
x = LSTM(64)(x)
output_reg = Dense(1, name='regression')(x)
output_clf = Dense(1, activation='sigmoid', name='classification')(x)

model = Model(inputs=input_layer, outputs=[output_reg, output_clf])
model.compile(
    optimizer='adam',
    loss={'regression': 'mse', 'classification': 'binary_crossentropy'},
    metrics={'regression': 'mae', 'classification': 'accuracy'}
)

model.summary()

es = EarlyStopping(patience=5, restore_best_weights=True)
model.fit(
    X_train, {'regression': y_train_reg, 'classification': y_train_clf},
    validation_split=0.2,
    epochs=50,
    batch_size=32,
    callbacks=[es]
)
```

Model: "functional"


Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 60, 20)	0	-
lstm (LSTM)	(None, 60, 128)	76,288	input_layer[0][0]
lstm_1 (LSTM)	(None, 64)	49,408	lstm[0][0]
regression (Dense)	(None, 1)	65	lstm_1[0][0]
classification (Dense)	(None, 1)	65	lstm_1[0][0]

Total params: 125,826 (491.51 KB)


Trainable params: 125,826 (491.51 KB)

Non-trainable params: 0 (0.00 B)


Epoch 1/50

16/16  14s 164ms/step - classification_accuracy: 0.5096 - classification_loss: 0.7013 - loss: 0.8008 - regression_loss: 0.0995 - regression_mae: 0.2399 - val_classification_accuracy: 0.5520 - val_classification_loss: 0.6900 - val_loss: 0.7071 - val_regression_loss: 0.0173 - val_regression_mae: 0.1217


Epoch 2/50

16/16  4s 129ms/step - classification_accuracy: 0.5472 - classification_loss: 0.6882 - loss: 0.6958 - regression_loss: 0.0075 - regression_mae: 0.0706 - val_classification_accuracy: 0.4480 - val_classification_loss: 0.6951 - val_loss: 0.7026 - val_regression_loss: 0.0070 - val_regression_mae: 0.0707


Epoch 3/50

16/16  2s 122ms/step - classification_accuracy: 0.5605 - classification_loss: 0.6817 - loss: 0.6866 - regression_loss: 0.0051 - regression_mae: 0.0567 - val_classification_accuracy: 0.5520 - val_classification_loss: 0.6870 - val_loss: 0.6908 - val_regression_loss: 0.0050 - val_regression_mae: 0.0575


Epoch 4/50

16/16  2s 114ms/step - classification_accuracy: 0.5246 - classification_loss: 0.6901 - loss: 0.6955 - regression_loss: 0.0052 - regression_mae: 0.0579 - val_classification_accuracy: 0.5440 - val_classification_loss: 0.6812 - val_loss: 0.6834 - val_regression_loss: 0.0026 - val_regression_mae: 0.0397


Epoch 5/50

16/16  2s 112ms/step - classification_accuracy: 0.6102 - classification_loss: 0.6626 - loss: 0.6683 - regression_loss: 0.0057 - regression_mae: 0.0573 - val_classification_accuracy: 0.5520 - val_classification_loss: 0.7190 - val_loss: 0.7196 - val_regression_loss: 0.0028 - val_regression_mae: 0.0429


Epoch 6/50

16/16  3s 142ms/step - classification_accuracy: 0.5700 - classification_loss: 0.6713 - loss: 0.6800 - regression_loss: 0.0086 - regression_mae: 0.0757 - val_classification_accuracy: 0.4560 - val_classification_loss: 0.6893 - val_loss: 0.6926 - val_regression_loss: 0.0023 - val_regression_mae: 0.0369


Epoch 7/50

16/16  3s 157ms/step - classification_accuracy: 0.5878 - classification_loss: 0.6609 - loss: 0.6651 - regression_loss: 0.0042 - regression_mae: 0.0512 - val_classification_accuracy: 0.5440 - val_classification_loss: 0.6699 - val_loss: 0.6737 - val_regression_loss: 0.0042 - val_regression_mae: 0.0533


Epoch 8/50

16/16  2s 124ms/step - classification_accuracy: 0.6861 - classification_loss: 0.6205 - loss: 0.6273 - regression_loss: 0.0067 - regression_mae: 0.0668 - val_classification_accuracy: 0.5840 - val_classification_loss: 0.6788 - val_loss: 0.6827 - val_regression_loss: 0.0037 - val_regression_mae: 0.0490


Epoch 9/50

16/16  2s 114ms/step - classification_accuracy: 0.7097 - classification_loss: 0.5916 - loss: 0.5959 - regression_loss: 0.0041 - regression_mae: 0.0475 - val_classification_accuracy: 0.5920 - val_classification_loss: 0.7235 - val_loss: 0.7270 - val_regression_loss: 0.0035 - val_regression_mae: 0.0481


Epoch 10/50

16/16  3s 114ms/step - classification_accuracy: 0.7283 - classification_loss: 0.5557 - loss: 0.5598 - regression_loss: 0.0043 - regression_mae: 0.0513 - val_classification_accuracy: 0.6560 - val_classification_loss: 0.6578 - val_loss: 0.6654 - val_regression_loss: 0.0068 - val_regression_mae: 0.0695


Epoch 11/50

16/16  2s 114ms/step - classification_accuracy: 0.7283 - classification_loss: 0.5642 - loss: 0.5679 - regression_loss: 0.0039 - regression_mae: 0.0471 - val_classification_accuracy: 0.5760 - val_classification_loss: 0.6783 - val_loss: 0.6869 - val_regression_loss: 0.0094 - val_regression_mae: 0.0898


Epoch 12/50

16/16  3s 209ms/step - classification_accuracy: 0.7120 - classification_loss: 0.5624 - loss: 0.5724 - regression_loss: 0.0102 - regression_mae: 0.0862 - val_classification_accuracy: 0.7280 - val_classification_loss: 0.6117 - val_loss: 0.6150 - val_regression_loss: 0.0029 - val_regression_mae: 0.0413


Epoch 13/50

16/16  2s 114ms/step - classification_accuracy: 0.7597 - classification_loss: 0.5276 - loss: 0.5338 - regression_loss: 0.0064 - regression_mae: 0.0645 - val_classification_accuracy: 0.7120 - val_classification_loss: 0.6162 - val_loss: 0.6204 - val_regression_loss: 0.0040 - val_regression_mae: 0.0524


Epoch 14/50

16/16  3s 115ms/step - classification_accuracy: 0.7548 - classification_loss: 0.5003 - loss: 0.5045 - regression_loss: 0.0038 - regression_mae: 0.0494 - val_classification_accuracy: 0.6160 - val_classification_loss: 0.6670 - val_loss: 0.6698 - val_regression_loss: 0.0028 - val_regression_mae: 0.0429

Epoch 15/50

16/16  3s 127ms/step - classification_accuracy: 0.7348 - classification_loss: 0.5382 - loss: 0.5416 - regression_loss: 0.0033 - regression_mae: 0.0461 - val_classification_accuracy: 0.7120 - val_classification_loss: 0.5571 - val_loss: 0.5597 - val_regression_loss: 0.0015 - val_regression_mae: 0.0298

Epoch 16/50

16/16  2s 115ms/step - classification_accuracy: 0.7961 - classification_loss: 0.4520 - loss: 0.4546 - regression_loss: 0.0027 - regression_mae: 0.0418 - val_classification_accuracy: 0.7280 - val_classification_loss: 0.4546 - val_loss: 0.4572 - val_regression_loss: 0.0015 - val_regression_mae: 0.0298

tion_loss: 0.5659 - val_loss: 0.5699 - val_regression_loss: 0.0034 - val_regression_mae: 0.0451
Epoch 17/50
16/16 ————— 3s 162ms/step - classification_accuracy: 0.7955 - classification_loss: 0.4296 - loss: 0.4330 - regression_loss: 0.0033 - regression_mae: 0.0426 - val_classification_accuracy: 0.7600 - val_classification_loss: 0.5319 - val_loss: 0.5368 - val_regression_loss: 0.0027 - val_regression_mae: 0.0382
Epoch 18/50
16/16 ————— 3s 164ms/step - classification_accuracy: 0.7552 - classification_loss: 0.4941 - loss: 0.5000 - regression_loss: 0.0057 - regression_mae: 0.0603 - val_classification_accuracy: 0.7360 - val_classification_loss: 0.5200 - val_loss: 0.5229 - val_regression_loss: 0.0020 - val_regression_mae: 0.0347
Epoch 19/50
16/16 ————— 2s 117ms/step - classification_accuracy: 0.8369 - classification_loss: 0.3931 - loss: 0.3988 - regression_loss: 0.0055 - regression_mae: 0.0580 - val_classification_accuracy: 0.7680 - val_classification_loss: 0.5104 - val_loss: 0.5161 - val_regression_loss: 0.0036 - val_regression_mae: 0.0518
Epoch 20/50
16/16 ————— 3s 115ms/step - classification_accuracy: 0.8316 - classification_loss: 0.3664 - loss: 0.3683 - regression_loss: 0.0030 - regression_mae: 0.0416 - val_classification_accuracy: 0.7920 - val_classification_loss: 0.4788 - val_loss: 0.4858 - val_regression_loss: 0.0051 - val_regression_mae: 0.0618
Epoch 21/50
16/16 ————— 2s 118ms/step - classification_accuracy: 0.8730 - classification_loss: 0.3268 - loss: 0.3298 - regression_loss: 0.0030 - regression_mae: 0.0449 - val_classification_accuracy: 0.7440 - val_classification_loss: 0.4634 - val_loss: 0.4671 - val_regression_loss: 0.0020 - val_regression_mae: 0.0314
Epoch 22/50
16/16 ————— 2s 118ms/step - classification_accuracy: 0.8650 - classification_loss: 0.3296 - loss: 0.3334 - regression_loss: 0.0038 - regression_mae: 0.0489 - val_classification_accuracy: 0.7280 - val_classification_loss: 0.4763 - val_loss: 0.4822 - val_regression_loss: 0.0031 - val_regression_mae: 0.0418
Epoch 23/50
16/16 ————— 3s 183ms/step - classification_accuracy: 0.8412 - classification_loss: 0.3569 - loss: 0.3615 - regression_loss: 0.0040 - regression_mae: 0.0501 - val_classification_accuracy: 0.8000 - val_classification_loss: 0.4643 - val_loss: 0.4687 - val_regression_loss: 0.0029 - val_regression_mae: 0.0432
Epoch 24/50
16/16 ————— 2s 132ms/step - classification_accuracy: 0.8949 - classification_loss: 0.2955 - loss: 0.2986 - regression_loss: 0.0035 - regression_mae: 0.0466 - val_classification_accuracy: 0.7680 - val_classification_loss: 0.4909 - val_loss: 0.5006 - val_regression_loss: 0.0070 - val_regression_mae: 0.0736
Epoch 25/50
16/16 ————— 2s 117ms/step - classification_accuracy: 0.8499 - classification_loss: 0.3112 - loss: 0.3154 - regression_loss: 0.0041 - regression_mae: 0.0514 - val_classification_accuracy: 0.7680 - val_classification_loss: 0.4427 - val_loss: 0.4462 - val_regression_loss: 0.0034 - val_regression_mae: 0.0435
Epoch 26/50
16/16 ————— 2s 117ms/step - classification_accuracy: 0.8677 - classification_loss: 0.3085 - loss: 0.3115 - regression_loss: 0.0026 - regression_mae: 0.0398 - val_classification_accuracy: 0.8000 - val_classification_loss: 0.4544 - val_loss: 0.4568 - val_regression_loss: 0.0016 - val_regression_mae: 0.0302
Epoch 27/50
16/16 ————— 3s 117ms/step - classification_accuracy: 0.8637 - classification_loss: 0.2921 - loss: 0.2943 - regression_loss: 0.0020 - regression_mae: 0.0349 - val_classification_accuracy: 0.8160 - val_classification_loss: 0.4979 - val_loss: 0.4999 - val_regression_loss: 0.0026 - val_regression_mae: 0.0381
Epoch 28/50
16/16 ————— 3s 148ms/step - classification_accuracy: 0.8973 - classification_loss: 0.2762 - loss: 0.2781 - regression_loss: 0.0025 - regression_mae: 0.0399 - val_classification_accuracy: 0.7760 - val_classification_loss: 0.4983 - val_loss: 0.5030 - val_regression_loss: 0.0018 - val_regression_mae: 0.0317
Epoch 29/50
16/16 ————— 3s 168ms/step - classification_accuracy: 0.8626 - classification_loss: 0.2943 - loss: 0.2959 - regression_loss: 0.0019 - regression_mae: 0.0344 - val_classification_accuracy: 0.7840 - val_classification_loss: 0.4813 - val_loss: 0.4853 - val_regression_loss: 0.0020 - val_regression_mae: 0.0340
Epoch 30/50
16/16 ————— 2s 115ms/step - classification_accuracy: 0.8965 - classification_loss: 0.2828 - loss: 0.2852 - regression_loss: 0.0026 - regression_mae: 0.0412 - val_classification_accuracy: 0.7680 - val_classification_loss: 0.4794 - val_loss: 0.4861 - val_regression_loss: 0.0048 - val_regression_mae: 0.0603

Out[14]: <keras.src.callbacks.history.History at 0x7a93dc7a0f50>

```
In [15]: y_pred_reg, y_pred_clf = model.predict(X_test)
y_pred_reg_2d = np.concatenate([y_pred_reg, np.zeros_like(y_pred_reg)], axis=1)
y_pred_reg_inv = scaler_y.inverse_transform(y_pred_reg_2d[:, 0])

y_test_reg_2d = np.concatenate([y_test_reg, np.zeros_like(y_test_reg)], axis=1)
y_test_reg_inv = scaler_y.inverse_transform(y_test_reg_2d[:, 0])
```

9/9 ————— 1s 76ms/step

```
In [16]: from sklearn.metrics import mean_absolute_error, median_absolute_error, r2_score, mean_absolute_percentage_error

mae = mean_absolute_error(y_test_reg_inv, y_pred_reg_inv)
medae = median_absolute_error(y_test_reg_inv, y_pred_reg_inv)
mape = np.mean(np.abs((y_test_reg_inv - y_pred_reg_inv) / y_test_reg_inv)) * 100
r2 = r2_score(y_test_reg_inv, y_pred_reg_inv)

print('Regression Evaluation')
print(f'MAE : {mae:.4f}')
print(f'MedAE : {medae:.4f}')
print(f'MAPE : {mape:.2f}%')
print(f'R^2 : {r2:.4f}')
print('-----')

fig, axs = plt.subplots(1, 2, figsize=(16, 6))
```



```

axs[0].plot(y_test_reg_inv[:360], label='Actual', linewidth=2)
axs[0].plot(y_pred_reg_inv[:360], label='Predicted', linestyle='--')
axs[0].set_title('BTC Price Prediction (Line Chart)')
axs[0].set_xlabel('Time Step')
axs[0].set_ylabel('BTC Price (USD)')
axs[0].legend()
axs[0].grid(True)

axs[1].scatter(y_test_reg_inv, y_pred_reg_inv, alpha=0.5)
axs[1].plot([min(y_test_reg_inv), max(y_test_reg_inv)],
            [min(y_test_reg_inv), max(y_test_reg_inv)], '--', label='Ideal')
axs[1].set_title('Actual vs Predicted (Scatter Plot)')
axs[1].set_xlabel('Actual Price')
axs[1].set_ylabel('Predicted Price')
axs[1].legend()
axs[1].grid(True)

plt.tight_layout()
plt.show()

```

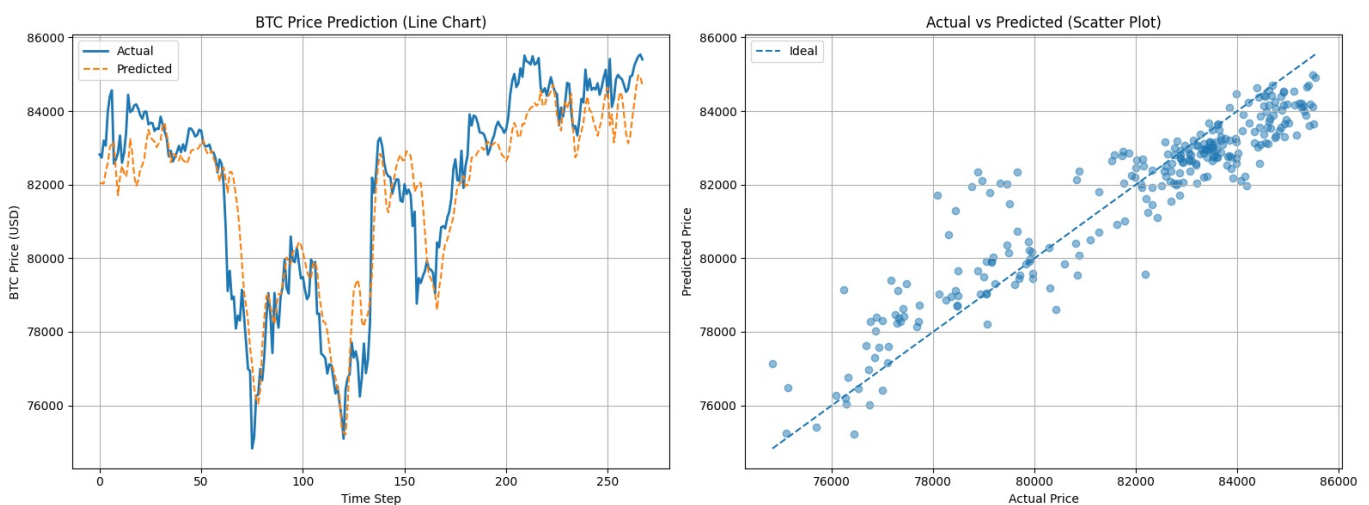
Regression Evaluation

MAE : 814.0456

MedAE : 652.1306

MAPE: 1.00%

R² : 0.8538



Wait...

Is the model actually good or is there a problem? I'm always skeptical of the models I build, let's do a walk forward test.

```

In [17]: walk_preds = []
walk_truth = []

input_seq = X_test[0].copy()

for i in range(len(X_test)):
    pred_reg, pred_clf = model.predict(input_seq[np.newaxis, :, :], verbose=0)
    pred = pred_reg[0][0]
    walk_preds.append(pred)
    walk_truth.append(y_test_reg[i][0])

    if i + 1 < len(X_test):
        next_row = X_test[i + 1, -1, :]
        next_row_with_pred = next_row.copy()
        next_row_with_pred[0] = pred
        input_seq = np.vstack([input_seq[1:], [next_row_with_pred]])

```

```

In [18]: walk_preds_2d = np.concatenate([np.array(walk_preds).reshape(-1, 1), np.zeros_like(np.array(walk_preds).reshape(-1, 1))])
walk_preds_inv = scaler_y.inverse_transform(walk_preds_2d[:, 0])
walk_truth_2d = np.concatenate([np.array(walk_truth).reshape(-1, 1), np.zeros_like(np.array(walk_truth).reshape(-1, 1))])
walk_truth_inv = scaler_y.inverse_transform(walk_truth_2d[:, 0])


mae_walk = mean_absolute_error(walk_truth_inv, walk_preds_inv)
mape_walk = np.mean(np.abs((walk_truth_inv - walk_preds_inv) / walk_truth_inv)) * 100
r2_walk = r2_score(walk_truth_inv, walk_preds_inv)

print(f"[Walk-Forward] MAE : {mae_walk:.2f}")
print(f"[Walk-Forward] MAPE : {mape_walk:.2f}%")
print(f"[Walk-Forward] R2 : {r2_walk:.4f}")

```

[Walk-Forward] MAE : 926.88
[Walk-Forward] MAPE : 1.14%
[Walk-Forward] R² : 0.8172

```
In [19]: plt.figure(figsize=(12, 6))
plt.plot(walk_truth_inv[:300], label='Actual', linewidth=2)
plt.plot(walk_preds_inv[:300], label='Predicted (Walk)', linestyle='--')
plt.title('Walk-Forward Prediction vs Actual (Last 300 Candle)')
plt.xlabel('Step')
plt.ylabel('BTC Price (USD)')
plt.legend()
plt.grid(True)
plt.show()
```



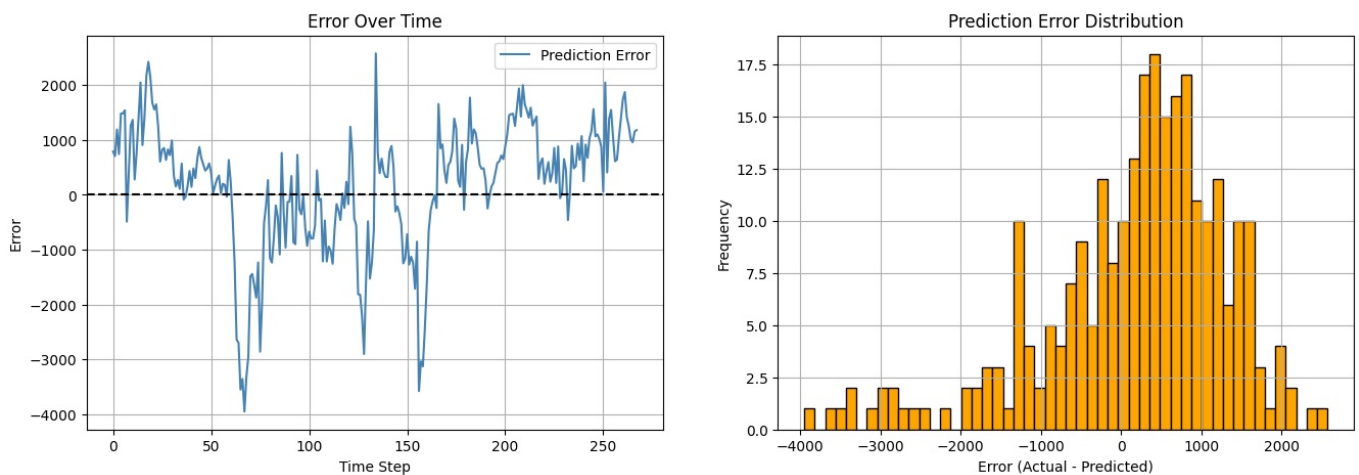
```
In [20]: error = walk_truth_inv.flatten() - walk_preds_inv.flatten()

fig, axs = plt.subplots(1, 2, figsize=(16, 5))

axs[0].plot(error, label='Prediction Error', color='steelblue')
axs[0].axhline(0, color='black', linestyle='--')
axs[0].set_title("Error Over Time")
axs[0].set_xlabel("Time Step")
axs[0].set_ylabel("Error")
axs[0].legend()
axs[0].grid(True)


axs[1].hist(error, bins=50, color='orange', edgecolor='black')
axs[1].set_title("Prediction Error Distribution")
axs[1].set_xlabel("Error (Actual - Predicted)")
axs[1].set_ylabel("Frequency")
axs[1].grid(True)

plt.show()
```



```
In [21]: n = 100
plt.figure(figsize=(12, 5))
```

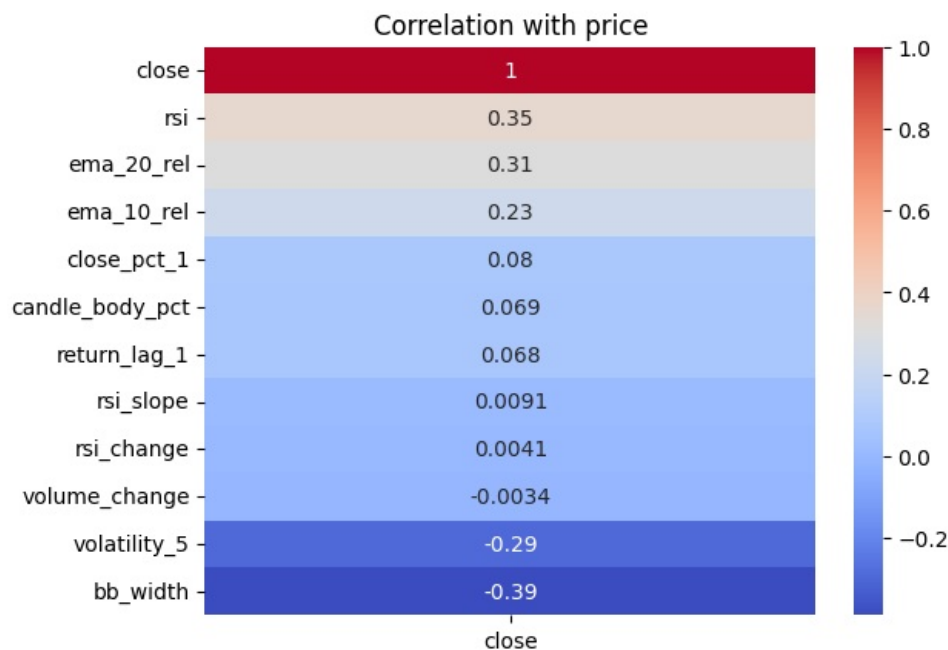
```
plt.plot(walk_truth_inv[-n:], label="Actual")
plt.plot(walk_preds_inv[-n:], label="Predicted", linestyle="--")
plt.title("Zoom: Predicted vs Actual (Last 100 Candles)")
plt.grid(True)
plt.legend()
plt.show()
```



In [22]: `import seaborn as sns`

```
selected_features = ['candle_body_pct', 'bb_width', 'close_pct_1', 'return_lag_1', 'volatility_5', 'volume_change', 'rsi_slope', 'rsi_change', 'volume_change', 'volatility_5', 'bb_width']
target_correlation = df_analyze[selected_features].corr()

sns.heatmap(target_correlation[['close']].sort_values(by='close', ascending=False), annot=True, cmap='coolwarm')
plt.title("Correlation with price")
plt.show()
```



Forecasting

This part is more interesting than the previous part, because we can earn money if our prediction is correct hehehe. Kidding hehehe. Actually I don't want to trust the model because there are too many simplifications here, and I know quant trade has more data and more complicated models to understand bitcoin.

So this is like a disclaimer to understand the context, this part doesn't need to be fully believed because what I'm doing here is just to train my thinking and how I can apply the knowledge I've gained so far in Data Science. Just for learning.

```
In [23]: forecast = []
input_seq = X_test[-1].copy()

for _ in range(12):
    pred_reg, pred_clf = model.predict(input_seq[np.newaxis, :, :], verbose=0)
    pred = pred_reg[0][0]
```

```

forecast.append(pred)

new_row = np.zeros_like(input_seq[0])
new_row[0] = pred
new_row[1:] = input_seq[-1, 1:]

input_seq = np.vstack([input_seq[1:], new_row])

forecast_scaled = np.array(forecast).reshape(-1, 1)
forecast_scaled_2d = np.concatenate([forecast_scaled, np.zeros_like(forecast_scaled)], axis=1)
forecast_actual = scaler_y.inverse_transform(forecast_scaled_2d)[: , 0].flatten()

```

```

In [24]: print("BTC price forecast (within the next 12 candles):")
forecast_actual = pd.DataFrame(forecast_actual, columns=['Price'])

forecast_actual.index = [i for i in range(1, 13)]

print(forecast_actual)


historical = df_analyze['close'].iloc[-60:].values
combined = np.concatenate([historical, forecast_actual['Price'].values])

plt.figure(figsize=(14, 6))
plt.plot(combined, label='BTC Price (Past + Forecast)', color='blue')
plt.axvline(x=60, color='gray', label='Forecast Start')
plt.title('BTC Price: Last 10 Days + Next 12 Hours Forecast')
plt.xlabel('Index')
plt.ylabel('BTC Price (USD)')
plt.legend()
plt.grid(True)
plt.show()

```

BTC price forecast (within the next 12 candles):

	Price
1	84711.046875
2	84638.000000
3	84742.328125
4	84868.296875
5	84985.921875
6	85106.773438
7	85169.953125
8	85141.023438
9	85043.062500
10	84919.750000
11	84809.750000
12	84735.085938



Sentiment Analysis

In the world of trading, there are actually two different groups, one is people who just analyze the price (technical analysis), the other is people who analyze the macro-micro economy to understand what will happen to the price (fundamental analysis). And I want to combine both of them into one group here (to help me decide what to do with the current price, whether buy, sell or hold).

Actually I wanted to get the data from twitter but I couldn't get it, so I will get from two sources, first from CoinDesk and second from CryptoPanic

```

In [25]: import requests

```

```

from bs4 import BeautifulSoup

def scrape_coindesk_with_content(max_articles=10):
    base_url = "https://www.coindesk.com/"
    response = requests.get(base_url)
    soup = BeautifulSoup(response.content, 'html.parser')

    articles = []
    seen_urls = set()

    for link_tag in soup.find_all('a', href=True):
        href = link_tag['href']
        full_url = href if href.startswith("http") else f"https://www.coindesk.com{href}"
        if "/" in href and "video" not in href and full_url not in seen_urls and href.startswith('/markets', '.'):
            seen_urls.add(full_url)
            try:
                article_page = requests.get(full_url)
                article_soup = BeautifulSoup(article_page.content, 'html.parser')

                title_tag = article_soup.find(['h1', 'h2'])
                title = title_tag.get_text(strip=True) if title_tag else "No title"
                paragraphs = article_soup.find_all('p')
                content = " ".join(p.get_text(strip=True) for p in paragraphs[:5])

                articles.append({
                    'url': full_url,
                    'title': title,
                    'content': content
                })

                if len(articles) >= max_articles:
                    break

            except Exception as e:
                print(f"Error fetching {full_url}: {e}")
                continue

    return articles

news_articles = scrape_coindesk_with_content(25)
news_articles[0:3]

```

```

Out[25]: [{'url': 'https://www.coindesk.com/markets/2025/04/15/dogecoin-slumps-3-bitcoin-steady-around-usd85k-as-traders-fear-u-s-recession',
  'title': 'Dogecoin Slumps 3%, Bitcoin Steady Around $85K as Traders Fear U.S. Recession',
  'content': "Dogecoin (DOGE) shed 3% while bitcoin (BTC) and ether (ETH) remained flat in the past 24 hours as tariff concerns gradually subsided among traders, though fears of a U.S. recession increased in betting markets. 'Prominent financial figures have started to warn that the U.S. is heading into an imminent recession, with betting markets placing 40% to 60% odds of one happening in 2025,' Augustine Fan, head of insights at SignalPlus, told CoinDesk in a Telegram message. 'Our view is that it probably doesn't matter, as sentiment often frames reality, not the other way around.' 'As such, crypto has benefited from the recent shake-out, as equities have been realizing higher volatility than Bitcoin through the risk-off move. A beggar-thy-neighbour policy with tariffs has pushed spot gold to ATHs, with BTC finally regaining some of its long-lost 'store of value' narrative,' Fan added. Crypto majors tracked by the broad-based CoinDesk 20 (CD20) slid nearly 2%, data shows, with DOGE leading losses. Solana's SOL, tron (TRX) and Cardano's ADA lost as much as 2.5%, BNB Chain's BNB and xrp (XRP) were little changed as bitcoin clung to the $85,000 level. Mantra's OM token showed a 20% rise over the past 24 hours to trade at 63 cents in Asian morning hours Tuesday, following a bizarre sell-off that saw it lose 90% within an hour late Sunday. A recovery plan is in the works, its CEO said in an interview following the plunge, though market watchers remain sceptical of any promises.'"},
  {'url': 'https://www.coindesk.com/business/2025/04/14/nomura-s-laser-digital-denies-involvement-in-mantra-crash',
  'title': "Nomura's Laser Digital Denies Involvement in Mantra Crash",
  'content': "Switzerland-based trading firm Laser Digital, which is part of the Nomura Group, has denied any involvement in the Mantra token flash crash that saw OM lose 90% of its value. 'Assertions circulating on social media that link Laser to 'investor selling' are factually incorrect and misleading,' the firm wrote on X. Laser Digital went on to share its controlled Mantra wallet addresses, none of which show deposits to exchanges or selling activity. Speculation remains rife over why OM collapsed so violently. The Mantra team insist it was due to wider market pressures and centralized exchanges forcibly closing positions, which led to a liquidation cascade. OKX stated that the price volatility occurred due to a spike in trading volume coupled with an initial price decline across various exchanges outside of OKX, before spreading to the wider market.'"},
  {'url': 'https://www.coindesk.com/markets/2025/04/14/bitcoin-facing-cloud-resistance-at-usd85k-neutralized-risk-reward-for-bulls',
  'title': "Bitcoin Faces 'Cloud Resistance' at $85K, Neutralizes Risk-Reward for Bulls: Godbole",
  'content': "This is a daily technical analysis by CoinDesk analyst and Chartered Market Technician Omkar Godbole. In markets, securing the best entry point is often half the battle, as timing and level significantly influence success by skewing the risk-reward ratio in traders' favor. While bitcoin's (BTC) near-term outlook may appear constructive with increased demand for bullish bets in the options market, the cryptocurrency's proximity to key resistance that capped the upside in recent months means the risk-reward profile for those looking to capitalize on the bullish prospects is less favorable. Since Saturday, BTC has been pushing against the lower boundary of the 'Ichimoku cloud' at around $85K. Developed by a Japanese journalist in the 1960s, the Ichimoku cloud is a technical analysis indicator that offers a comprehensive view of market momentum, support, and resistance levels. The indicator comprises five lines: Leading Span A, Leading Span B, Conversion Line or Tenkan-Sen (T), Base Line or Kijun-Sen (K) and a lagging closing price line.'"}]

```

```

In [26]: api_key = "api_key"

base_url = "https://cryptopanic.com/api/v1/posts/"

params = {
    'auth_token': api_key,
    'filter': 'trending',
    'currencies': 'BTC',
    'regions': 'en',
    'kind': 'news',
    'public': True
}

response = requests.get(base_url, params=params)
news_data = []

if response.status_code == 200:
    data = response.json()
    for item in data['results']:
        title = item['title']
        url = item['url']
        source = item['source']['title']
        created_at = item['created_at']

        if 'currencies' in item:
            currencies = ', '.join([c['code'] for c in item['currencies']])
        else:
            currencies = ''

        news_item = {
            'judul': title,
            'url': url,
            'sumber': source,
            'waktu': created_at,
            'mata_uang': currencies,
            'content': title
        }

        news_data.append(news_item)

news_data[0:3]

```

```

Out[26]: [{'judul': 'MicroStrategy Acquires 3,459 Bitcoins for $285.8 Million, Total Holdings Reach 531,644 BTC',
  'url': 'https://cryptopanic.com/news/21072568/MicroStrategy-Acquires-3459-Bitcoins-for-2858-Million-Total-Ho-
  ldings-Reach-531644-BTC?mtm_campaign=API-OFA',
  'sumber': 'DeFi News',
  'waktu': '2025-04-14T14:30:40Z',
  'mata_uang': 'BTC',
  'content': 'MicroStrategy Acquires 3,459 Bitcoins for $285.8 Million, Total Holdings Reach 531,644 BTC'},
 {'judul': 'Saylor's Strategy adds 3,459 Bitcoin, now holds 531,644 BTC',
  'url': 'https://cryptopanic.com/news/21069129/Saylors-Strategy-adds-3459-Bitcoin-now-holds-531644-BTC?mtm_cam-
  paign=API-OFA',
  'sumber': 'CryptoBriefing',
  'waktu': '2025-04-14T12:03:51Z',
  'mata_uang': 'BTC',
  'content': 'Saylor's Strategy adds 3,459 Bitcoin, now holds 531,644 BTC'},
 {'judul': 'Saylor signals new Bitcoin buy after Strategy reports nearly $6 billion Q1 unrealized loss',
  'url': 'https://cryptopanic.com/news/21056879/Saylor-signals-new-Bitcoin-buy-after-Strategy-reports-nearly-6-
  billion-Q1-unrealized-loss?mtm_campaign=API-OFA',
  'sumber': 'CryptoBriefing',
  'waktu': '2025-04-13T14:44:50Z',
  'mata_uang': 'BTC',
  'content': 'Saylor signals new Bitcoin buy after Strategy reports nearly $6 billion Q1 unrealized loss'}]

```

```

In [27]: import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import pandas as pd

nltk.download('vader_lexicon')

sia = SentimentIntensityAnalyzer()

all_articles = news_articles + news_data
df_news = pd.DataFrame(all_articles)
df_news['content'] = df_news['content'].astype(str)

def vader_sentiment(text):
    score = sia.polarity_scores(text)['compound']
    if score >= 0.05:
        return 'Positive'
    elif score <= -0.05:
        return 'Negative'

```

```

        else:
            return 'Neutral'

df_news['sentimen'] = df_news['content'].apply(vader_sentiment)

df_news[['title', 'sentimen']].head()

```

```

[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
[nltk_data]   Package vader_lexicon is already up-to-date!

```

```

Out[27]:

```

	title	sentimen
0	Dogecoin Slumps 3%, Bitcoin Steady Around \$85K...	Negative
1	Nomura's Laser Digital Denies Involvement in M...	Negative
2	Bitcoin Faces 'Cloud Resistance' at \$85K, Neut...	Positive
3	XRP, SOL and ADA Flash Bullish Patterns as Tra...	Positive
4	Metaplanet Becomes 10th Largest Public Bitcoin...	Positive

```

In [28]: sentiment_counts = df_news['sentimen'].value_counts()
dominant_sentiment = sentiment_counts.idxmax()
jumlah_dominan = sentiment_counts.max()
total_berita = sentiment_counts.sum()

print("News Sentiment Conclusion (VADER)")
print(f"Total news analyzed: {total_berita}")
print(f"Dominant sentiment: {dominant_sentiment} ({jumlah_dominan} news)")
print("\nComplete distribution:")
print(sentiment_counts)

```

```

News Sentiment Conclusion (VADER)
Total news analyzed: 37
Dominant sentiment: Positive (14 news)

```

```

Complete distribution:
sentimen
Positive    14
Negative    13
Neutral     10
Name: count, dtype: int64

```

Decision

```

In [29]: prediction_sentiment = ''
if df_analyze.iloc[-1]['close'] > forecast_actual['Price'].mean():
    prediction_sentiment = 'Negative'
elif df_analyze.iloc[-1]['close'] < forecast_actual['Price'].mean():
    prediction_sentiment = 'Positive'
else:
    prediction_sentiment = 'Neutral'

def signal(prediction_sentiment, dominant_sentiment):
    if prediction_sentiment == "Positive":
        if dominant_sentiment == "Positive":
            return "BUY"
        elif dominant_sentiment == "Neutral":
            return "WAIT"
        else:
            return "CAUTION"
    elif prediction_sentiment == "Negative":
        if dominant_sentiment == "Negative":
            return "SELL"
        elif dominant_sentiment == "Positive":
            return "CAUTION"
        else:
            return "WAIT"
    else:
        return "HOLD"

trading_signal = signal(prediction_sentiment, dominant_sentiment)

```

```

In [30]: from matplotlib.patches import FancyBboxPatch

color_map = {
    "BUY": "#27ae60",
    "SELL": "#c0392b",
    "WAIT": "#f39c12",
    "CAUTION": "#f1c40f",
    "HOLD": "#7f8c8d"
}

```



```
message_map = {
    "BUY": "Time to buy!",
    "SELL": "Consider selling",
    "WAIT": "Wait for confirmation",
    "CAUTION": "Caution, the market is unstable",
    "HOLD": "Hold the position"
}

fig, ax = plt.subplots(figsize=(6, 3))
ax.set_xlim(0, 1)
ax.set_ylim(0, 1)
ax.axis('off')

box = FancyBboxPatch((0.05, 0.2), 0.9, 0.6,
                     boxstyle="round,pad=0.05",
                     linewidth=2,
                     facecolor=color_map.get(trading_signal, "#7f8c8d"),
                     edgecolor='black')

ax.add_patch(box)

ax.text(0.5, 0.6, f"Signal: {trading_signal}", fontsize=18, weight='bold',
        color='white', ha='center', va='center')

ax.text(0.5, 0.4, message_map.get(trading_signal, ""), fontsize=12,
        color='white', ha='center', va='center')

plt.title("Today Recommendations", fontsize=14)
plt.show()
```

Today Recommendations

Signal: CAUTION

Caution, the market is unstable